

statLearn R LAB.6(b) RIDGE REGRESSION AND LASSO

First, need to install and load the 'glmnet' package, glmnet does not use model formula language, so we must set up an 'x' and 'y'

```
library(ISLR)
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 3.3.2
```

```
## Loading required package: Matrix
```

```
## Loading required package: foreach
```

```
## Loaded glmnet 2.0-10
```

```
Hitters=na.omit(Hitters)
with(Hitters,sum(is.na(Salary)))
```


```
## [1] 0
```

```
x = model.matrix(Salary~.-1, data=Hitters)
y = Hitters$Salary
```

Ridge Regression: Shrinkage of Coefficients


To fit Ridge-Regression model we calling 'glmnet' with 'alpha=0' (see helpfile; to fit Lasso model, alpha is set to 1). 'cv.glmnet' function will do cross-validation for us

```
fit.ridge = glmnet(x, y, alpha=0)
plot(fit.ridge, xvar="lambda", label=TRUE)
```




Coefficients are plotted as a function of log lambda; RR models are penalized by the SumSquares of Coefficients controlled by parameter lambda: $RSS + \lambda \cdot \text{penalty}$ Coefficients shrink towards zero, as lambda increases

```
cv.ridge = cv.glmnet(x, y, alpha=0)
plot(cv.ridge)
```




The Lasso: does Shrinkage and Variable Selection Now, we fit lasso model; using the default 'alpha=1'
Lasso is similar to RR, but has slightly different penalty
uses absolute value of coefficients, shrinking some to zero

```
fit.lasso = glmnet(x, y, alpha=1)
plot(fit.lasso, xvar="lambda", label=TRUE)
```



```
plot(fit.lasso, xvar="dev", label=TRUE)
```

```
coef(cv.lasso)
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) 115.3773590
## AtBat      .
## Hits       1.4753071
## HmRun      .
## Runs       .
## RBI        .
## Walks      1.6566947
## Years      .
## CAtBat     .
## CHits      .
## CHmRun     .
## CRuns      0.1660465
## CRBI       0.3453397
## CWalks     .
## LeagueA    .
## LeagueN    .
## DivisionW  -19.2435216
## PutOuts    0.1000068
## Assists    .
## Errors     .
## NewLeagueN .
```

Use earlier TRAIN/VALIDATION to select ‘lambda’ for the lasso: It tries to fit 100 values of lambda, and if nothing changes, then it stops. Use best lambda and fit to whole data set

```
set.seed(1)
train = sample(seq(263), 180, replace=FALSE)
lasso.tr = glmnet(x[train,], y[train])
lasso.tr
```

```
##
## Call:  glmnet(x = x[train, ], y = y[train])
##
##           Df      %Dev   Lambda
## [1,]    0 0.00000 246.40000
## [2,]    1 0.05013 224.50000
## [3,]    1 0.09175 204.60000
## [4,]    2 0.13840 186.40000
## [5,]    2 0.18000 169.80000
## [6,]    3 0.21570 154.80000
## [7,]    3 0.24710 141.00000
## [8,]    3 0.27320 128.50000
## [9,]    4 0.30010 117.10000
## [10,]   4 0.32360 106.70000
## [11,]   4 0.34310  97.19000
## [12,]   4 0.35920  88.56000
## [13,]   5 0.37360  80.69000
```


##	[14,]	5	0.38900	73.52000
##	[15,]	5	0.40190	66.99000
##	[16,]	5	0.41260	61.04000
##	[17,]	5	0.42140	55.62000
##	[18,]	5	0.42880	50.67000
##	[19,]	5	0.43490	46.17000
##	[20,]	5	0.43990	42.07000
##	[21,]	5	0.44410	38.33000
##	[22,]	5	0.44760	34.93000
##	[23,]	6	0.45140	31.83000
##	[24,]	7	0.45480	29.00000
##	[25,]	7	0.45770	26.42000
##	[26,]	7	0.46010	24.07000
##	[27,]	8	0.46220	21.94000
##	[28,]	8	0.46380	19.99000
##	[29,]	8	0.46520	18.21000
##	[30,]	8	0.46630	16.59000
##	[31,]	8	0.46730	15.12000
##	[32,]	8	0.46810	13.78000
##	[33,]	9	0.47110	12.55000
##	[34,]	9	0.47380	11.44000
##	[35,]	9	0.47620	10.42000
##	[36,]	10	0.48050	9.49500
##	[37,]	9	0.48450	8.65200
##	[38,]	10	0.48770	7.88300
##	[39,]	10	0.49360	7.18300
##	[40,]	11	0.49890	6.54500
##	[41,]	12	0.50450	5.96300
##	[42,]	12	0.51010	5.43400
##	[43,]	13	0.51470	4.95100
##	[44,]	13	0.51850	4.51100
##	[45,]	13	0.52170	4.11000
##	[46,]	14	0.52440	3.74500
##	[47,]	14	0.52670	3.41200
##	[48,]	15	0.52870	3.10900
##	[49,]	15	0.53030	2.83300
##	[50,]	15	0.53160	2.58100
##	[51,]	16	0.53280	2.35200
##	[52,]	17	0.53420	2.14300
##	[53,]	18	0.53580	1.95300
##	[54,]	18	0.53760	1.77900
##	[55,]	18	0.53890	1.62100
##	[56,]	18	0.54000	1.47700
##	[57,]	18	0.54090	1.34600
##	[58,]	18	0.54160	1.22600
##	[59,]	18	0.54220	1.11700
##	[60,]	18	0.54280	1.01800
##	[61,]	18	0.54320	0.92770
##	[62,]	18	0.54360	0.84530
##	[63,]	18	0.54380	0.77020


```
## [64,] 19 0.54410 0.70180
## [65,] 19 0.54430 0.63940
## [66,] 19 0.54450 0.58260
## [67,] 19 0.54470 0.53090
## [68,] 19 0.54490 0.48370
## [69,] 20 0.54510 0.44070
## [70,] 20 0.54520 0.40160
## [71,] 20 0.54530 0.36590
## [72,] 20 0.54540 0.33340
## [73,] 20 0.54550 0.30380
## [74,] 20 0.54560 0.27680
## [75,] 20 0.54570 0.25220
## [76,] 20 0.54570 0.22980
## [77,] 20 0.54580 0.20940
## [78,] 20 0.54580 0.19080
## [79,] 20 0.54590 0.17380
## [80,] 20 0.54590 0.15840
## [81,] 20 0.54590 0.14430
## [82,] 20 0.54590 0.13150
## [83,] 20 0.54600 0.11980
## [84,] 19 0.54600 0.10920
## [85,] 19 0.54600 0.09948
## [86,] 19 0.54600 0.09064
## [87,] 19 0.54600 0.08259
## [88,] 20 0.54600 0.07525
## [89,] 20 0.54600 0.06856
```

```
pred = predict(lasso.tr, x[-train,])
dim(pred)
```

```
## [1] 83 89
```

```
rmse = sqrt(apply((y[-train]-pred)^2, 2, mean))
plot(log(lasso.tr$lambda), rmse, type="b", xlab="Log(lambda)")
```



```
lam.best = lasso.tr$lambda[order(rmse)[1]]  
lam.best
```

```
## [1] 19.98706
```

```
coef(lasso.tr, s=lam.best)
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              1
## (Intercept) 107.9416686
## AtBat      .
## Hits       0.1591252
## HmRun      .
## Runs       .
## RBI        1.7340039
## Walks      3.4657091
## Years      .
## CAtBat     .
## CHits      .
## CHmRun     .
## CRuns      0.5386855
## CRBI       .
## CWalks     .
## LeagueA    -30.0493021
## LeagueN    .
## DivisionW  -113.8317016
## PutOuts    0.2915409
## Assists    .
## Errors     .
## NewLeagueN 2.0367518
```

=====